

Lee Yan Cheng A0199141B

GitHub link: <https://github.com/yanchenglee98/OTOT-A2-A3>

Instructions:

- 1) install metrics server: `kubectl apply -f https://github.com/kubernetes-sigs/metrics-server/releases/latest/download/components.yaml`
- 2) manually edit Deployment manifest to ensure tls is working:

`kubectl -nkube-system edit deploy/metrics-server`

add a flag `--kubelet-insecure-tls` to `deployment.spec.containers[].args[]`

restart the Deployment using `kubectl -nkube-system rollout restart deploy/metrics-server`

- 3) apply hpa: `kubectl apply -f ./demo/a3/hpa.yaml`
- 4) spam refresh at `http://localhost/app/`
- 5) check hpa event log: `kubectl describe hpa`
- 6) apply zone aware: `kubectl apply -f ./demo/a3/deployment-zone-aware.yaml`
- 7) check zone aware allocation: `kubectl get po -lapp=a2-zone-aware -owide`

Screenshot proof:

Before scale:

```
C:\Users\yanch\Downloads\cs3219\OTOT-A2-A3>kubectl describe hpa
Warning: autoscaling/v2beta2 HorizontalPodAutoscaler is deprecated in v1.23+, unavailable in v1.26+; use autoscaling/v2 HorizontalPodAutoscaler
Name: a2
Namespace: default
Labels: <none>
Annotations: <none>
CreationTimestamp: Thu, 20 Oct 2022 21:20:53 +0800
Reference: Deployment/a2
Metrics: ( current / target )
  resource cpu on pods (as a percentage of request): 45% (9m) / 50%
Min replicas: 1
Max replicas: 10
Deployment pods: 1 current / 1 desired
Conditions:
  Type            Status  Reason                        Message
  ----            -
  AbleToScale     True    ReadyForNewScale             recommended size matches current size
  ScalingActive   True    ValidMetricFound             the HPA was able to successfully calculate a replica count from cpu resource utilization (percentage of request)
  ScalingLimited  False   DesiredWithinRange           the desired count is within the acceptable range
Events:
```

After scale:

```
C:\Users\yanch\Downloads\cs3219\OTOT-A2-A3>kubectl describe hpa
Warning: autoscaling/v2beta2 HorizontalPodAutoscaler is deprecated in v1.23+, unavailable in v1.26+; use autoscaling/v2 HorizontalPodAutoscaler
Name: a2
Namespace: default
Labels: <none>
Annotations: <none>
CreationTimestamp: Thu, 20 Oct 2022 21:20:53 +0800
Reference: Deployment/a2
Metrics: ( current / target )
  resource cpu on pods (as a percentage of request): 16% (2m) / 50%
Min replicas: 1
Max replicas: 10
Deployment pods: 10 current / 10 desired
Conditions:
  Type            Status  Reason                        Message
  ----            -
  AbleToScale     True    ScaleDownStabilized          recent recommendations were higher than current one, applying the highest recent recommendation
  ScalingActive   True    ValidMetricFound             the HPA was able to successfully calculate a replica count from cpu resource utilization (percentage of request)
  ScalingLimited  True    TooManyReplicas              the desired replica count is more than the maximum replica count
Events:
  Type    Reason              Age    From                      Message
  ----    -
  Normal  SuccessfulRescale   7m39s  horizontal-pod-autoscaler  New size: 1; reason: All metrics below target
  Normal  SuccessfulRescale   5m54s  horizontal-pod-autoscaler  New size: 5; reason: cpu resource utilization (percentage of request) above target
  Normal  SuccessfulRescale   5m9s   horizontal-pod-autoscaler  New size: 9; reason: cpu resource utilization (percentage of request) above target
  Normal  SuccessfulRescale   4m38s  horizontal-pod-autoscaler  New size: 10; reason: cpu resource utilization (percentage of request) above target
```

HPA event output:

Events:				
Type	Reason	Age	From	Message
----	-----	----	----	-----
Normal	SuccessfulRescale	7m39s	horizontal-pod-autoscaler	New size: 1; reason: All metrics below target
Normal	SuccessfulRescale	5m54s	horizontal-pod-autoscaler	New size: 5; reason: cpu resource utilization (percentage of request) above target
Normal	SuccessfulRescale	5m9s	horizontal-pod-autoscaler	New size: 9; reason: cpu resource utilization (percentage of request) above target
Normal	SuccessfulRescale	4m38s	horizontal-pod-autoscaler	New size: 10; reason: cpu resource utilization (percentage of request) above target